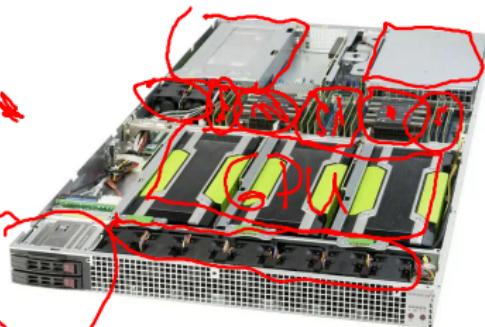


A Primer on Computers and Data Centers

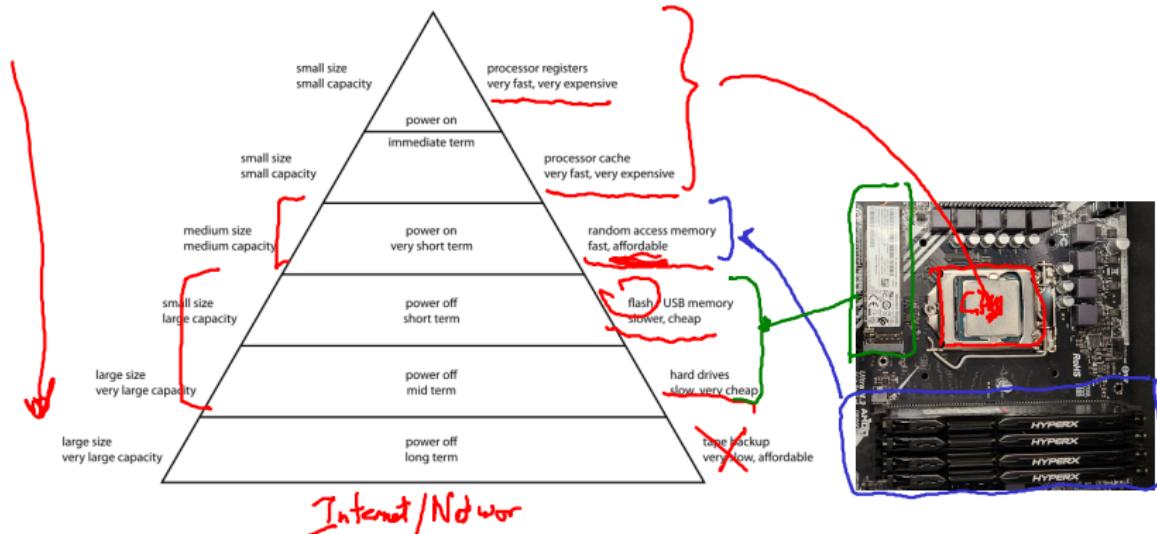


- ▶ Dual socket CPU: 96+ cores each
- ▶ Memory: 128GB to 1TB
- ▶ Storage: 500GB SSD plus multiple SSDs/HDDs for data storage

This is only about 5-20 times as powerful as your computer ...

The Computer Memory Hierarchy

Computer Memory Hierarchy



Key point:

- ▶ To be very fast, the 'inner loop' must fit inside cache
- ▶ To be fast enough, the 'inner loop' must fit inside RAM

Big Data

Problem: There are some really big datasets out there . . .

Competitions With Largest Datasets

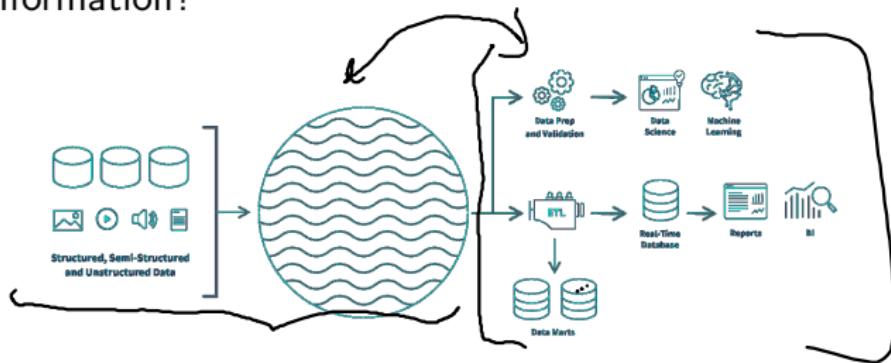
| | Report | Script | Input | Output | Logs | Comments (1) |
|---|-------------|--|-------|--------|------|--------------|
| | Competition | | | | | Size |
| → | 1 | Diabetic Retinopathy Detection | | | | |
| | 2 | American Epilepsy Society Seizure Prediction Challenge | | | | |
| | 3 | Avito Duplicate Ads Detection | | | | |
| | 4 | Microsoft Malware Classification Challenge (BIG 2015) | | | | |
| | 5 | GE Flight Quest | | | | |
| | 6 | Draper Satellite Image Chronology | | | | |
| | 7 | CHALEARN Gesture Challenge | | | | |
| | 8 | Second Annual Data Science Bowl | | | | |
| | 9 | Flight Quest 2: Flight Optimization, Main Phase | | | | |
| | 10 | Belkin Energy Disaggregation Competition | | | | |

Another page claims there are 526 datasets on Kaggle > 100GB

Data Analytics with Spark

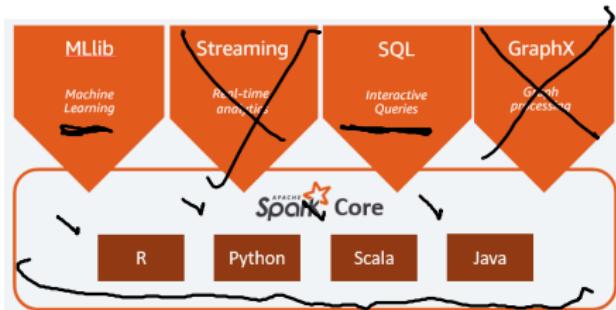
"Data analytics converts raw data into actionable insights."

Problem: How to convert an organization's huge data sets into useful information?



- ▶ "Apache Spark is an open-source, distributed processing system used for big data workloads."
- ▶ "It utilizes in-memory caching, and optimized query execution for fast analytic queries against data of any size."
- ▶ "Spark provides an interface for programming clusters with implicit data parallelism and fault tolerance."

Apache Spark Internals



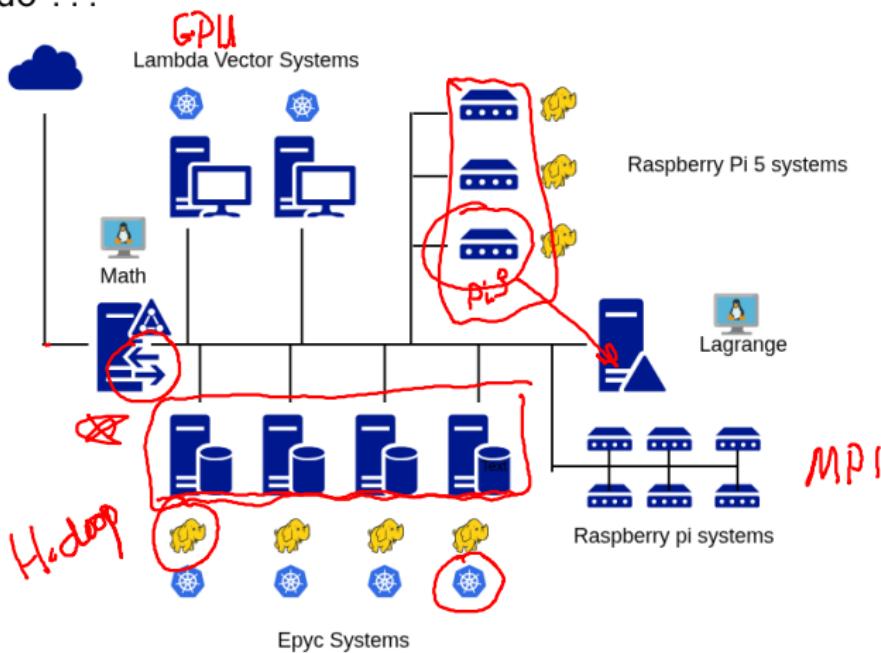
You

- ▶ Spark is built on Hadoop Mapreduce, “a programming model . . . for processing and generating big data sets with a parallel and distributed algorithm on a cluster.”
- ▶ Spark “has its architectural foundation in . . . a read-only multiset of data items distributed over a cluster of machines, that is maintained in a fault-tolerant way.” [Dataframes]
- ▶ Spark uses HDFS, which “provides the scalable, fault-tolerant storage layer, while Spark acts as the high-speed, in-memory data processing engine that operates on the data stored in HDFS.”

Apache Spark at Knox

Question: Why have you never used Apache Spark? Because you don't have a cluster ...

but we do ...



Apache Spark at Knox



Spark Master at spark://lagrange:7077

URL: spark://lagrange:7077

REST URL: spark://lagrange:6066 (cluster mode)

Workers: 4 Alive, 0 Dead, 0 Decommissioned, 0 Unknown

Cores in use: 256 Total, 9 Used

Memory in use: 499.0 GiB Total, 15.0 GiB Used

Resources in use:

Applications: 1 Running, 4 Completed

Drivers: 0 Running (0 Waiting), 0 Completed (0 Killed, 0 Failed, 0 Error, 0 Relaunching)

Status: ALIVE (Environment, Log)

Workers (4)

| Worker Id | Address | State | Cores | Memory |
|--|------------------|-------|-------------|--------------------------|
| worker-20260209195803-10.90.1.25-32997 | 10.90.1.25:32997 | ALIVE | 64 (0 Used) | 124.8 GiB (0.0 B Used) |
| worker-20260212122337-10.90.1.24-42373 | 10.90.1.24:42373 | ALIVE | 64 (3 Used) | 124.8 GiB (5.0 GiB Used) |
| worker-20260212123631-10.90.1.23-42885 | 10.90.1.23:42885 | ALIVE | 64 (3 Used) | 124.8 GiB (5.0 GiB Used) |
| worker-20260212124218-10.90.1.37-44471 | 10.90.1.37:44471 | ALIVE | 64 (3 Used) | 124.8 GiB (5.0 GiB Used) |

Running Applications (1)

| Application ID | Name | Cores | Memory per Executor | Resources Per Executor | Submitted Time | User | State | Duration |
|--------------------------------|-----------|-------|---------------------|------------------------|---------------------|--------|---------|----------|
| app-20260213094345-0004 (kill) | ClassDemo | 9 | 5.0 GiB | | 2026/02/13 09:43:45 | aleahy | RUNNING | 30.7 h |

This cluster uses an 8TB HDFS filesystem with double redundancy on each file block spread over the same four systems

Using Spark: Basic Concepts

- ▶ A Spark cluster can be accessed with a *SparkSession* from several programming languages (Python, R, Scala, ...)

A screenshot of a Jupyter Notebook interface. The code cell contains the following Python code:

```
[30]: from pyspark.sql import SparkSession  
spark = SparkSession.builder \  
    .appName("ClassDemo") \  
    .getOrCreate()
```

The word "builder" in the second line is circled in red. A curly brace on the left side of the slide groups this point with the next one.

- ▶ Spark typically uses the *parquet* file format

A diagram illustrating the internal structure of a Parquet file. It shows a grid of data with the following columns:

| | Column 1 | Column 2 | Column 3 | Column 4 | Column 5 |
|--------------|--------------|---------------|----------|------------|----------|
| Product | Bell T-Shirt | John Doe | USA | 2023-01-01 | 100 |
| Customer | John Doe | Maria Adams | UK | 2023-01-02 | 200 |
| Country | | Antonio Grant | USA | 2023-01-03 | 300 |
| Date | | Maria Adams | UK | 2023-01-04 | 100 |
| Sales Amount | | John Doe | USA | 2023-01-05 | 500 |
| | | | | | 200 |

The data is organized into three Row Groups, each highlighted with a different color (red, orange, green). Within each row group, individual columns are grouped into Column Chunks. An arrow points from the text "Column chunks in one row group" to the boundary between the first and second column of the first row group. A curly brace on the left side of the slide groups this point with the next one.

ID, sex, number,

{

- ▶ Parquet is a *columnar* file format, which allows *pruning*
- ▶ Columnar data can be compressed, because it is similar
- ▶ Parquet supports *predicate pushdown*, in which a *where* filter “filters the data in the database query”
- ▶ Spark supports all this, which speeds up queries up to 10-100x

The Spark API

Key point: Spark has to solve the same problems as Pandas does
... See the Spark API reference (for 4.1.0)

Another key point: The Spark API has been evolving, so there are multiple entrypoints—some better than others

Some Entrypoints for the API:

- ▶ Spark SQL - data filtering and aggregating
- ▶ Pandas API - data input/output and transformation
- ▶ MLlib - many familiar machine learning models

See the Spark API Reference:

<https://spark.apache.org/docs/latest/api/python/reference/index.html>

... or try your favorite AI for example syntax ...