

Maximum Likelihood Estimators

(Raschke)

The MLE for logistic regression is:

$$\mathcal{L}(\mathbf{w}, b | \mathbf{x}) = p(y | \mathbf{x}; \mathbf{w}, b) = \prod_{i=1}^n p(y^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}, b) = \prod_{i=1}^n (\sigma(z^{(i)}))^{y^{(i)}} (1 - \sigma(z^{(i)}))^{1-y^{(i)}}$$

In practice, it is easier to maximize the (natural) log of this equation, which is called the **log-likelihood** function:

$$l(\mathbf{w}, b | \mathbf{x}) = \log \mathcal{L}(\mathbf{w}, b | \mathbf{x}) = \sum_{i=1}^n [y^{(i)} \log(\sigma(z^{(i)})) + (1 - y^{(i)}) \log(1 - \sigma(z^{(i)}))]$$

The formula for the gradient in $w = w - \eta \nabla L(w, b)$:

$$\frac{\partial L}{\partial w_j} = \underbrace{\frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial w_j}}_{\text{Apply chain rule}} \quad \text{where } a = \sigma(z) = \frac{1}{1 + e^{-z}}$$

1) Derive terms separately:

2) Combine via chain rule and simplify:

$$\begin{aligned} \frac{\partial L}{\partial a} &= \frac{a - y}{a - a^2} \\ \frac{da}{dz} &= \frac{e^{-z}}{(1 + e^{-z})^2} = a \cdot (1 - a) \\ \frac{\partial z}{\partial w_j} &= x_j \end{aligned}$$

] ——————> $\frac{\partial L}{\partial z} = a - y$] ——————> $\frac{\partial L}{\partial w_j} = (a - y)x_j = -(y - a)x_j$

The Logistic Binary Classifier

Suppose X is a function that gives outputs 0, 1. The **logistic function** is the function

$$p(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}. \quad (4.2)$$

We can rewrite this into an **odds formula**:

$$\frac{p(X)}{1 - p(X)} = e^{\beta_0 + \beta_1 X}. \quad (4.3)$$

and take the logarithm to get the **logit** function:

$$\log\left(\frac{p(X)}{1 - p(X)}\right) = \beta_0 + \beta_1 X. \quad (4.4)$$

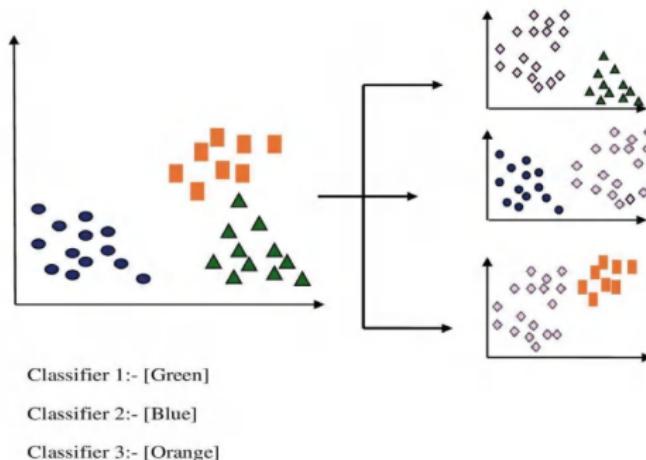
We can extend this using multivariate regression as well . . .

One vs. the Rest

Question: How to extend this to K (more than two) classes?

There are a couple of ways ...

In a **one-vs-rest classifier** (or **OvR** or **OvA**), we train a single binary classifiers several times, each time comparing one of the K categories to the remainder—thus generating K classifiers. Then choose the category K which has the highest probability.



Multinomial Logistic Regression

Idea: That approach works for *all* binary classifiers. There is another approach for logistic regression.

Choose a **baseline** class (say the K th class). We want a function that computes

$$\Pr(Y = k|X = x) = \frac{e^{\beta_{k0} + \beta_{k1}x_1 + \dots + \beta_{kp}x_p}}{1 + \sum_{l=1}^{K-1} e^{\beta_{l0} + \beta_{l1}x_1 + \dots + \beta_{lp}x_p}} \quad (4.10)$$

for $k = 1, \dots, K-1$, and

$$\Pr(Y = K|X = x) = \frac{1}{1 + \sum_{l=1}^{K-1} e^{\beta_{l0} + \beta_{l1}x_1 + \dots + \beta_{lp}x_p}}. \quad (4.11)$$

It is not hard to show that for $k = 1, \dots, K-1$,

$$\log \left(\frac{\Pr(Y = k|X = x)}{\Pr(Y = K|X = x)} \right) = \beta_{k0} + \beta_{k1}x_1 + \dots + \beta_{kp}x_p. \quad (4.12)$$

As in OvR, the idea is to come up with a formula for $P(Y = k|X = x)$ and then take the maximum over these values.

Finding $P(Y = k | X = x)$

We have:

$$\ln \frac{\Pr(Y_i = k)}{\Pr(Y_i = K)} = \beta_k \cdot \mathbf{X}_i, \quad 1 \leq k < K.$$

If we exponentiate both sides and solve for the probabilities, we get:

$$\Pr(Y_i = k) = \Pr(Y_i = K) e^{\beta_k \cdot \mathbf{X}_i}, \quad 1 \leq k < K$$

Using the fact that all K of the probabilities must sum to one, we find:

$$\begin{aligned} \Pr(Y_i = K) &= 1 - \sum_{j=1}^{K-1} \Pr(Y_i = j) \\ &= 1 - \sum_{j=1}^{K-1} \Pr(Y_i = K) e^{\beta_j \cdot \mathbf{X}_i} \Rightarrow \Pr(Y_i = K) \\ &= \frac{1}{1 + \sum_{j=1}^{K-1} e^{\beta_j \cdot \mathbf{X}_i}}. \end{aligned}$$

We can use this to find the other probabilities:

$$\Pr(Y_i = k) = \frac{e^{\beta_k \cdot \mathbf{X}_i}}{1 + \sum_{j=1}^{K-1} e^{\beta_j \cdot \mathbf{X}_i}}, \quad 1 \leq k < K.$$

SKLearn Logistic Regression

LogisticRegression

```
class
sklearn.linear_model.LogisticRegression(penalty='deprecated',
*, C=1.0, l1_ratio=0.0, dual=False, tol=0.0001,
fit_intercept=True, intercept_scaling=1, class_weight=None,
random_state=None, solver='lbfgs', max_iter=100, verbose=0,
warm_start=False, n_jobs=None) [source]
```

The solvers 'lbfgs', 'newton-cg', 'newton-cholesky' and 'sag' support only L2 regularization with primal formulation, or no regularization. The 'liblinear' solver supports both L1 and L2 regularization (but not both, i.e. elastic-net), with a dual formulation only for the L2 penalty. The Elastic-Net (combination of L1 and L2) regularization is only supported by the 'saga' solver.

For [multiclass](#) problems (whenever `n_classes >= 3`), all solvers except 'liblinear' optimize the (penalized) multinomial loss. 'liblinear' only handles binary classification but can be extended to handle multiclass by using [`OneVsRestClassifier`](#).